



MECHATRONICS




MATHWARE

ASSEMBLY



ELECTRONICS



SOFTWARE



Ranges



Ranges

- Why do we need ranges?
 - STL algorithms
 - Legacy Iterators vs C++20 iterators
- What is the definition of a range?
- Views and adaptors
- Sentinels

Why do we need ranges?

Iterators are fundamentally unsafe

There are no built-in safety checks or sentinels to prevent iterator use after invalidation

```
std::vector<int> v{1, 2, 3};  
auto it = v.begin();  
v.push_back(4);    // may reallocate, invalidates iterators  
int x = *it;       // undefined behaviour
```

Iterators and sentinels

Iterators come in pairs of the same type: begin and end.

This makes some patterns awkward (e.g., searching until a terminator).

```
// Can't stop at '\0' without computing std::end(...)
for (auto it = std::begin(arr); it != std::end(arr); ++it) {
    ...
}
```

Upside: You could roll your own iterator

Downside: Complexity

Composability

Iterators do not compose well.

You typically end up writing loops instead of 'chaining' operations.

```
std::vector<int> result;  
for (auto it = v.begin(); it != v.end(); ++it) {  
    if (!predicate(*it)) continue;  
    auto transformed = transform(*it);  
    if (stop_condition(transformed)) break;  
    result.push_back(transformed);  
}
```

Upside: Memory usage

Downside: Readability

Granted: you could use 'intermediate containers' but you incur a memory and cpu loss.
Hence do not compose 'well'.

```
std::vector<int> intermediate;  
std::copy_if(v.begin(), v.end(), std::back_inserter(intermediate),  
            [](int x) { return x % 2 == 0; });  
  
std::vector<int> result;  
std::transform(intermediate.begin(), intermediate.end(), std::back_inserter(result),  
              [](int x) { return x * 10; });
```

Upside: Readability

Downside: CPU usage, Memory usage

Eager vs lazy

Standard STL algorithms are **eager**. They process the entire range immediately, producing a full result before the next step begins.

- Large-scale datasets that don't fit in RAM.
- Streaming data, where the 'end iterator' is not reachable
 - Custom iterators

```
std::vector<int> intermediate;  
std::copy_if(v.begin(), v.end(), std::back_inserter(intermediate),  
            [](int x) { return x % 2 == 0; });  
  
std::vector<int> result;  
std::transform(intermediate.begin(), intermediate.end(), std::back_inserter(result),  
              [](int x) { return x * 10; });
```

How are we supposed to transform 1 element only exit?
say, early exit? Copy has been done already.

How to handle failure?

If `std::find` fails to find anything, you must manually compare with `end()`.

- Dereferencing the `end()` iterator is UB.
- Compiler does not help you here.

Example

```
auto it = std::find(v.begin(), v.end(), 99);  
int x = *it;    // BUG if element not found
```

Iterator Invalidation

Different containers have different invalidation rules.

Container	push_back invalidates?
<code>std::vector</code>	Possibly
<code>std::list</code>	Never

Interface to containers

Decouple algorithms from data structures -> a single generic function can process any container type.

```
std::vector<int> vec = {10, 20, 30, 40};  
auto it = std::find(vec.begin(), vec.end(), 30);  
if (it != vec.end()) { /* found */ }
```

```
std::list<int> lst = {10, 20, 30, 40};  
auto it = std::find(lst.begin(), lst.end(), 40);  
if (it != lst.end()) { /* found */ }
```

ex1.cpp:

Convert a raw for loop to an STL algorithm.

The goal:

(re)discover the pain of using iterators.

Legacy Iterators (>C++20)

An iterator is an object that points to a specific element within a data structure, functioning much like a smart cursor.

Since C++20, we have 'named requirements':

cppreference	standard
LegacyInputIterator	Cpp17InputIterator
LegacyIterator	Cpp17Iterator

The standard specifies list of requirements but they are not enforced or checked.

Conceptification of a LegacyIterator ([Cpp17Iterator](#)):

```
template<class I>
concept __LegacyIterator =
    requires(I i)
    {
        { *i } -> __Referenceable;
        { ++i } -> std::same_as<I&>;
        { *i++ } -> __Referenceable;
    } && std::copyable<I>;
```

Named requirement

Failure to do so may result in very complex compiler diagnostics.

Legacy algorithms

Lets take a look at 1 declaration of `std::find`

```
template<class InputIterator, class T = iterator_traits<InputIterator>::value_type>  
constexpr InputIterator find(InputIterator first, InputIterator last, const T& value);
```

If we then look-up `InputIterator`:

If an algorithm's template parameter is named `InputIterator`,
the template argument shall meet the Cpp17InputIterator requirements ([input.iterators]).

This is contract is NOT enforced by the compiler.

`cppreference`:

InputIt must meet the requirements of LegacyInputIterator.

5 Core traits

To be useable by legacy algorithms (e.g. `std::find`), types NEED to have:

- `value_type`: The type of the element the iterator points to.
- `difference_type`: A signed integer type that can represent the distance between two iterators.
- `reference`: The reference type of the element (usually `value_type&`).
- `pointer`: The pointer type of the element (usually `value_type*`).
- `iterator_category`: A tag (e.g., `std::forward_iterator_tag`, `std::random_access_iterator_tag`) that tells algorithms which operations the iterator supports, enabling compile-time optimizations (tag dispatching).

Deprecated since C++17:

```
class Iterator
```

Base class that provides the 5 core traits.

There are layers of legacy

```

template <typename Iterator>
struct iterator_traits {
    using difference_type    = typename Iterator::difference_type;
    using value_type         = typename Iterator::value_type;
    using pointer            = typename Iterator::pointer;
    using reference          = typename Iterator::reference;
    using iterator_category  = typename Iterator::iterator_category;
};

template <typename T>
struct iterator_traits<T*> {
    using difference_type    = std::ptrdiff_t;
    using value_type         = std::remove_cv_t<T>;
    using pointer            = T*;
    using reference          = T&;
    using iterator_category  = std::random_access_iterator_tag;
};

```

```

#include <iterator>
struct InvalidIterator {
    // Missing: using difference_type = std::ptrdiff_t;
    using value_type = int;           // Present
    using pointer = int*;            // Present
    using reference = int&;          // Present
    using iterator_category = std::forward_iterator_tag; // Present
}
int main() {
    static_assert(std::is_same_v<std::iterator_traits<InvalidIterator>::value_type, int>);
}

```

```

<source>:10:73: error: 'value_type' is not a member of 'std::iterator_traits<InvalidIterator>'
 10 |         static_assert(std::is_same_v<std::iterator_traits<InvalidIterator>::value_type, int>);
    |

```

```

class MyIntIterator {
public:
    using iterator_category = std::forward_iterator_tag;
    using difference_type    = std::ptrdiff_t;
    using value_type         = int;
    using pointer             = int*;
    using reference           = int&;
    explicit MyIntIterator(pointer ptr) : m_ptr(ptr) {}
    reference operator*() const { return *m_ptr; }
    pointer operator->() { return m_ptr; }
    MyIntIterator& operator++() {
        m_ptr++;
        return *this;
    }
    MyIntIterator operator++(int) {
        MyIntIterator tmp = *this;
        ++(*this);
        return tmp;
    }
    friend bool operator==(const MyIntIterator& a, const MyIntIterator& b) {
        return a.m_ptr == b.m_ptr;
    }
    friend bool operator!=(const MyIntIterator& a, const MyIntIterator& b) {
        return a.m_ptr != b.m_ptr;
    }
private:
    pointer m_ptr;
};

```

Tag dispatching

```
template <typename Iterator, typename Distance>
void my_advance(Iterator& it, Distance n) {
    using category = typename std::iterator_traits<Iterator>::iterator_category;
    detail::advance_impl(it, n, category{});
}
```

```

namespace detail {
    // Overload for Input and Forward Iterators: O(N) complexity
    template <typename Iterator, typename Distance>
    void advance_impl(Iterator& it, Distance n, std::forward_iterator_tag) {
        // Must step through one by one
        while (n > 0) {
            ++it;
            --n;
        }
    }
    // Overload for Bidirectional Iterators: O(N) complexity but supports negative n
    template <typename Iterator, typename Distance>
    void advance_impl(Iterator& it, Distance n, std::bidirectional_iterator_tag) {
        if (n > 0) {
            while (n-- > 0) ++it;
        } else {
            while (n++ < 0) --it;
        }
    }
    template <typename Iterator, typename Distance>
    void advance_impl(Iterator& it, Distance n, std::random_access_iterator_tag) {
        // Direct jump using pointer arithmetic
        it += n;
    }
}

```

```
#include <iostream>
#include <vector>
#include <string>
#include <iterator>
struct File { //Invisible to the user
    std::string name;
    std::string content;
};
struct Directory {
    std::vector<File> files;
    //public API
    TextIterator begin() const { return TextIterator(files.begin()); }
    TextIterator end() const { return TextIterator(files.end()); }
};
```

We want to loop over a set of files in a directory.

The concept of a 'File' is an implementation detail the user of our lib does not know of this.

```

class TextIterator {
private:
    std::vector<File>::const_iterator file_it;
public:
    using iterator_category = std::forward_iterator_tag;
    using value_type        = std::string;
    using difference_type    = std::ptrdiff_t;
    using pointer            = const std::string*;
    using reference          = const std::string&;
    explicit TextIterator(std::vector<File>::const_iterator it) : file_it(it) {}
    // intercept the dereference to return the string, not the File
    reference operator*() const { return file_it->content; }
    pointer operator->() const { return &(file_it->content); }
    TextIterator& operator++() {
        ++file_it;
        return *this;
    }
    TextIterator operator++(int) {
        TextIterator tmp = *this;
        ++file_it;
        return tmp;
    }
    friend bool operator==(const TextIterator& a, const TextIterator& b) {
        return a.file_it == b.file_it;
    }
    friend bool operator!=(const TextIterator& a, const TextIterator& b) {
        return a.file_it != b.file_it;
    }
};

```

```
int main() {
    Directory my_dir = {
        {
            {"file1.txt", "Hello from the first file."},
            {"file2.txt", "This is the second file."},
            {"file3.txt", "And here is the third!"}
        }
    };
    // The loop only sees std::strings. The File structs are invisible.
    for (const std::string& text : my_dir) {
        std::cout << text << "\n";
    }
    return 0;
}
```

ex2.cpp:

Write a custom legacy iterator for image data.

At this point you might be thinking? SFINAE? Non compiler enforced contract?

```
using iterator_category = std::forward_iterator_tag;  
using value_type         = std::string;  
using difference_type    = std::ptrdiff_t;  
using pointer            = const std::string*;  
using reference          = const std::string&;
```

Feels like something from c++17? Yes!

C++20 Iterators

C++20: Enter concepts and type constraining.

- differentiate the weak, 'legacy' unchecked named-requirements from the new, compiler-enforced `std::` concepts.

Input / Output Iterators

The most basic iterators. Support single-pass read (Input) or write (Output) operations. Reading or writing consumes the element (e.g., stream iterators)..

```
template< class I >
concept input_iterator =
    std::input_or_output_iterator<I> &&
    std::indirectly_readable<I> &&
    std::derived_from<ITER_CONCEPT<I>, std::input_iterator_tag>;
```

Forward iterators

Multi-pass forward iteration. safe to copy the iterator and iterate over the same range multiple times without consuming it.

```
template< class I >
    concept forward_iterator =
        std::input_iterator<I> &&
        std::derived_from</*ITER_CONCEPT*/<I>, std::forward_iterator_tag> &&
        std::incrementable<I> &&
        std::sentinel_for<I, I>;
```

Bidirectional Iterators

Forward guarantees, plus the ability to iterate backwards step-by-step (supports --it).

```
template< class I >
    concept bidirectional_iterator =
        std::forward_iterator<I> &&
        std::derived_from</*ITER_CONCEPT*/<I>, std::bidirectional_iterator_tag> &&
        requires(I i) {
            { --i } -> std::same_as<I&>;
            { i-- } -> std::same_as<I>;
        };

```

Random access iterators

Constant time $O(1)$ jumps ($it + n$), distance calculations, and relational comparisons ($it1 < it2$).

```
template< class I >
    concept random_access_iterator =
        std::bidirectional_iterator<I> &&
        std::derived_from</*ITER_CONCEPT*/<I>, std::random_access_iterator_tag> &&
        std::totally_ordered<I> &&
        std::sized_sentinel_for<I, I> &&
        requires(I i, const I j, const std::iter_difference_t<I> n) {
            { i += n } -> std::same_as<I&>;
            { j +  n } -> std::same_as<I>;
            { n +  j } -> std::same_as<I>;
            { i -= n } -> std::same_as<I&>;
            { j -  n } -> std::same_as<I>;
            { j[n]  } -> std::same_as<std::iter_reference_t<I>>;
        };
};
```

Contiguous Iterators

Random access and guarantees elements are adjacent in physical memory (e.g., `std::vector`, `std::array`).

```
template< class I >
    concept contiguous_iterator =
        std::random_access_iterator<I> &&
        std::derived_from</*ITER_CONCEPT*/<I>, std::contiguous_iterator_tag> &&
        std::is_lvalue_reference_v<std::iter_reference_t<I>> &&
        std::same_as<std::iter_value_t<I>,
                    std::remove_cvref_t<std::iter_reference_t<I>>> &&
        requires(const I& i) {
            { std::to_address(i) } ->
            std::same_as<std::add_pointer_t<std::iter_reference_t<I>>>;
        };
};
```

- C-API Interoperability; `memcpy`

You might say: whats the point of the tags. eg: `input_iterator_tag`?

- Two different iterator types might share an identical API (they both support `*it`, `++it`, and `==`)
 - one might be a single-pass stream iterator: input iterator (`std::istream_iterator`)
 - the other a multi-pass container iterator: forward iterator (`std::vector::iterator`).

By explicitly specifying using `iterator_concept = std::forward_iterator_tag`;, the author of `forward_list` is telling the compiler:

```
Even though my syntax looks like a basic incrementable type,  
I guarantee that copying it is safe and it satisfies multi-pass semantics.
```

The tag is needed

Input Iterator (`std::input_iterator_tag`): Supports ++, but is single-pass.

Forward Iterator (`std::forward_iterator_tag`): Supports ++, and is multi-pass.

Using c++20 Iterators

Constrain type I1 and I2 to random_access_iterator.
Similar to the tag dispatch example earlier.

```
template <std::random_access_iterator I1,  
          std::random_access_iterator I2>  
bool same_distance(I1 first1, I1 last1,  
                  I2 first2, I2 last2) {  
    // O(1) pointer arithmetic guaranteed  
    return (last1 - first1) == (last2 - first2);  
}
```

This contract is checked by the compiler!

When worlds collide:

```
struct ModernIter {
    // We only provide the strict minimum aliases required for C++20 Concepts
    using value_type = int;
    using difference_type = std::ptrdiff_t;
    using iterator_concept = std::input_iterator_tag;
    // INTENTIONALLY MISSING: // using pointer = int*; // using reference = int&; // using iterator_category = std::input_iterator_tag;
    int* ptr;
    int& operator*() const { return *ptr; }
    ModernIter& operator++() { ++ptr; return *this; }
    // C++20 allows post-increment to return void for input iterators.
    void operator++(int) { ++ptr; }
    bool operator==(const ModernIter& o) const = default;
};
static_assert(std::input_iterator<ModernIter>);\
void test() {
    int arr[] = {1, 2, 3};
    auto it2 = std::find(ModernIter{arr}, ModernIter{arr+3}, 2);
}
```

```
stl_algo.h:3863:13: error: no type named 'value_type' in 'struct std::iterator_traits<ModernIter>'
3863 |         using _ValT = typename iterator_traits<_InputIterator>::value_type;
```

ex3.cpp:

Write a $\geq$ C++20 iterator to analyze packet data.

Ranges

By **decree**:

1. A range is an **iterator** and a **sentinel** that designate the beginning and end of the computation
2. An **iterator** and a **count** that designate the beginning and the number of elements to which the computation is to be applied

A **range** is any object that can produce a begin and an end iterator.

```
template< class T >
concept range = requires( T& t ) {
    ranges::begin(t);
    ranges::end (t);
};
```

T E; ranges::begin(E)

- If T is an array type, ranges::begin(E) is expression-equivalent to t + 0.
- If auto(t.begin()) is a valid expression whose type models input_or_output_iterator, ranges::begin(E) is expression-equivalent to auto(t.begin()).
- If T is a class or enumeration type and auto(begin(t)) is a valid expression whose type models input_or_output_iterator

ranges::begin is a customization point -> avoid namespace poisoning due to ADL

```

#include <iostream>
#include <vector>
#include <iterator>
namespace vendor { // --- Third-Party Graphics Library ---
    struct Color { float r, g, b; };
    template <typename Target>
    void begin(Target& t) {
        std::cout << "vendor::begin -> Started rendering batch!\n";
    }
}
template <typename Container>
void process_elements(Container& c) {
    using std::begin; // "Two-Step" idiom (unrelated)
    auto it = begin(c);
    std::cout << "Processing elements...\n";
}
int main() {
    std::vector<vendor::Color> colors = {{1,0,0}, {0,1,0}};
    process_elements(colors);
}

```

```
<source>:14:20: error: call of overloaded 'begin(std::vector<vendor::Color>&)' is ambiguous  
14 |         auto it = begin(c);
```

Niebloid: customization point

Eric Niebler

[Ranges-v3](#). The reference implementation for c++ ranges
cppreference also uses this term 'customization point'

```

struct _Begin
{
private:
    template<typename _Tp>
        static constexpr bool
        _S_noexcept()
        {
            if constexpr (is_array_v<remove_reference_t<_Tp>>)
                return true;
            else if constexpr (__member_begin<_Tp>)
                return noexcept(_GLIBCXX_AUTO_CAST(std::declval<_Tp&>().begin()));
            else
                return noexcept(_GLIBCXX_AUTO_CAST(begin(std::declval<_Tp&>())));
        }
public:
    template<__maybe_borrowed_range _Tp>
        requires is_array_v<remove_reference_t<_Tp>> || __member_begin<_Tp>
        || __adl_begin<_Tp>
        [[nodiscard, __gnu__::__always_inline__]]
        constexpr auto
        operator()( _Tp&& __t) const noexcept(_S_noexcept<_Tp&>())
        {
            if constexpr (is_array_v<remove_reference_t<_Tp>>)
            {
                static_assert(is_lvalue_reference_v<_Tp>);
                return __t + 0;
            }
            else if constexpr (__member_begin<_Tp>)
                return __t.begin();
            else
                return begin(__t);
        }
};

```

What is a range?

A range is any object that exposes an iteration domain.

The important idea is that we describe the whole traversal with one object.

```
std::vector<int> v = {5, 1, 4, 2};  
std::ranges::sort(v);  
std::ranges::for_each(v, [](int x) { std::cout << x << ' '; });
```

The old style was:

```
std::sort(v.begin(), v.end());
```

The new style says:

```
std::ranges::sort(v);
```

What can a range express?

A range can represent a container, a subrange, a view, or a custom stream of values. This is the same shape of operation on a wide variety of sources.

```
std::vector<int> v = {1, 2, 3, 4, 5, 6};  
auto r = std::ranges::subrange(v.begin(), v.begin() + 4);  
for (int x : r) {  
    std::cout << x << ' ';  
}
```

We are no longer forced to think in terms of a pair of iterators that must match exactly.

A range can be adapted lazily

A view is a range that does not own its data. It composes cheaply.

```
std::vector<int> v = {1, 2, 3, 4, 5, 6};  
auto r1 = std::ranges::subrange(v);  
auto r2 = std::ranges::take_view(r1, 4);  
auto r3 = std::ranges::filter_view(r2,  
    [](int x) { return x % 2 == 0; });  
  
for (int x : r3) {  
    std::cout << x << ' '  
}  

```

The important point is not the type. The important point is that the range remains a view.

From pipeline to view

Once you see a range as a data source, the next step is to transform that source without copying it.

```
std::vector<int> v = {1, 2, 3, 4, 5, 6};

auto r = v
    | std::ranges::views::take(4)
    | std::ranges::views::filter([](int x) { return x % 2 == 0; });

for (int x : r) {
    std::cout << x << ' ';
}
```

The arrow is not magic. It is a range being passed through another range adaptor.

A range is constrained by concepts

Most of the useful checking happens at compile time.

That is why the range family is organized around concepts.

```
template< ranges::random_access_range R, class Comp = ranges::less,  
         class Proj = std::identity >  
requires std::sortable<ranges::iterator_t<R>, Comp, Proj>  
constexpr ranges::borrowed_iterator_t<R>  
    sort( R&& r, Comp comp = {}, Proj proj = {} );
```

We already know this concept more or less

```
template< class T >  
concept random_access_range =  
    ranges::bidirectional_range<T> &&  
    std::random_access_iterator<ranges::iterator_t<T>>;
```

The key observation is that this is not based on the older iterator traits family. Ranges depend on the new and improved iterator concepts.

Projection: identity

A projection is a function that extracts a key before the comparator runs. This removes boilerplate and makes the comparison more expressive.

```
struct Lad {  
    std::string name;  
    int age;  
};  
  
int main() {  
    std::vector<Lad> theLads = {"Erik", 77}, {"Bob", 33}, {"Charlie", 53};  
    std::ranges::sort(theLads,  
        [](const auto& a, const auto& b) { return a.age > b.age; });  
    std::ranges::sort(theLads, std::ranges::greater{}, &Lad::age);  
    std::ranges::sort(theLads, [](int a, int b) { return a > b; }, &Lad::age);  
}
```

Projection

- The projection runs before the comparator sees the elements.
- It can be a member pointer, a function pointer, or a lambda.
- Projections are called twice per comparison, once for each side.

Projection in one sentence

We transform the element into one comparable key, then compare those keys.

```
std::ranges::sort(theLads, std::ranges::greater{}, &Lad::age);
```

No need to ask the caller to write an explicit comparator on the whole object.

A sentinel can replace an end iterator

If a range ends on a special terminator, a sentinel can describe that end.

```
struct semicolon_sentinel {};  
  
bool operator==(const char* it, semicolon_sentinel)  
{  
    return *it == ';';  
}  
  
const char* text = "abc;def";  
auto r = std::ranges::subrange(  
    text,  
    semicolon_sentinel{}  
);
```

The iterator can stay lightweight while the end condition is expressed by a sentinel.

Why sentinels matter

Not every range ends with an iterator of the same type.

By separating the traversal end from the iterator type, we can model richer sources.

```
for (auto it = std::ranges::begin(r); it != std::ranges::end(r); ++it) {  
    std::cout << *it;  
}
```

The concept is the same: one object describing a sequence and one end condition that need not match the iterator type exactly.

What is a view?

A view is a lightweight range object.

It *may* own the data it is looking at.

It simply gives a new way to observe a source range.

```
std::vector<int> data{1, 2, 3, 4, 5, 6};  
  
auto first_three = std::ranges::subrange(data.begin(), data.begin() + 3);  
  
for (int x : first_three) {  
    std::cout << x << ' ';    // 1 2 3  
}
```

The vector still owns the elements.

The view only describes a window over them.

A view can be a window over the same data

Think of a view as a lens.

The underlying data remains in the vector, but the view decides which slice is visible.

```
std::vector<int> data{1, 2, 3, 4, 5, 6};

auto left = std::ranges::subrange(data.begin(), data.begin() + 3);
auto right = data | std::views::drop(3);

for (int x : left) std::cout << x << ' ';
for (int x : right) std::cout << x << ' ';
```

Both ranges observe the same vector.

No copy is made. The view only changes how we see the storage.

Views compose

The idea of views is composability:

- range → transform → filter → slice → consume
- do not copy the underlying data
- the view is copied, not the container

```
std::vector<int> v{1, 2, 3, 4, 5, 6};  
auto r1 = std::ranges::subrange(v);  
auto r2 = std::ranges::take_view(r1, 4);  
auto r3 = std::ranges::filter_view(r2, [](int x){ return x % 2 == 0; });  
  
for (int x : r3)  
    std::cout << x << ' ';
```

Every step wraps the previous range without materializing a new container.

Pipeline syntax

Use the pipe operator to build a range expression from left to right.

```
std::vector<int> v{1, 2, 3, 4, 5, 6};  
auto r = v  
    | std::ranges::views::take(4)  
    | std::ranges::views::filter([](int x) { return x % 2 == 0; });  
for (int x : r)  
    std::cout << x << ' ';
```

Adaptors: factories for new view instances

An adaptor is a range factory.

It takes one range and returns another range-like object, usually a lightweight view.

No intermediate container is created!!

The purpose is to keep the pipeline lazy and composable:

range → filter → transform → take → consume

```
std::vector<int> data{1, 2, 3, 4, 5, 6};

auto evens = data
    | std::views::filter([](int x) { return x % 2 == 0; });

for (int x : evens) {
    std::cout << x << ' ';    // 2 4 6
}
```

What an adaptor really does

An adaptor answers the question:

"Given a range, can I produce a new view with a different shape?"

Common examples:

- `std::views::filter`: keep elements satisfying a predicate
- `std::views::transform`: map each element to another value
- `std::views::take`: stop after N elements

The adaptor does not eagerly copy or sort anything.

It composes with the next adaptor and with the algorithm that consumes the range.

```
std::vector<int> data{1, 2, 3, 4, 5, 6};

auto result = data
    | std::views::filter([](int x) { return x % 2 == 0; })
    | std::views::transform([](int x) { return x * 10; })
    | std::views::take(2);

for (int x : result) {
    std::cout << x << ' ';    // 20 40
}
```

The important idea is that every step is a view, and each view remains cheap to store and compose.

```
int main() {  
    std::vector data{1, 2, 3, 4, 5, 6};  
    auto evens = data | std::views::filter([](int x) { return x % 2 == 0; });  
    std::cout << get_type_name<decltype(evens)>() << "\n\n";  
    return 0;  
}
```

```
std::ranges::filter_view<std::ranges::ref_view<std::vector<int, std::allocator<int> > >, main::{lambda(int)#1}>
```

Ranges are complicated types

What is the real type?

```
#include <ranges>
#include <vector>
#include <iostream>
std::string demangle(const char* name);
int main()
{
    std::vector<int> v = {1,2,3,4};
    auto is_even = [](int x){ return x % 2 == 0; };
    std::ranges::filter_view fv(v, is_even);
    std::cout << demangle(typeid(fv).name()) << std::endl;
}
```

```
std::string demangle(const char* name)
{
    int status = 0;
    std::unique_ptr<char, void(*)(void*)> res{
        abi::__cxa_demangle(name, nullptr, nullptr, &status),
        std::free
    };
    return (status == 0) ? res.get() : name;
}
```

```
std::ranges::filter_view<std::ranges::ref_view<std::vector<int, std::allocator<int> > >, main::{lambda(int)#1}>
```

- Filter predicate is part of the type? Good --> EBO

Only possible because of Class Template Argument Deduction (CTAD, C++17) .

Not possible as return value. rv must be fully qualified (or deduced).

```
template< ranges::input_range V,  
         std::indirect_unary_predicate<ranges::iterator_t<V>> Pred >  
    requires ranges::view<V> && std::is_object_v<Pred>  
class filter_view  
    : public ranges::view_interface<filter_view<V, Pred>>
```

std::ranges::ref_view

The adaptor wraps the range in a `ref_view` to avoid copying the vector.

The view stores a reference to the original object, not a duplicate of the container.

```
std::vector<int> v = {1, 2, 3, 4};  
auto is_even = [](int x){ return x % 2 == 0; };  
auto factory = std::views::filter(is_even);  
auto view = v | factory;  
  
static_assert(  
    std::same_as<  
        decltype(view),  
        std::ranges::filter_view<std::ranges::ref_view<decltype(v)>, decltype(is_even)>  
    >  
>);
```

```
std::ranges::filter_view<std::ranges::ref_view<vector<int>>, predicate>
```

The predicate is the view transformation.

The storage strategy is `ref_view` , so the vector is not copied.

std::ranges::views::all

`std::views::all` is a range adapter that decides how to keep the source alive.

- If the input is an lvalue: `std::ranges::ref_view<T>` .
- If the input is an rvalue: `std::ranges::owning_view<T>` .

```
int main()
{
    std::vector<int> v = {1, 2, 3};
    auto view1 = std::views::all(v); // l-value case: keep a reference to the existing vector
    static_assert(
        std::is_same_v<decltype(view1), std::ranges::ref_view<std::vector<int>>>,
        "Expected ref_view for lvalue"
    );
    auto view2 = std::views::all(std::vector<int>{1, 2, 3}); //r-value case: move the vector into an owning view
    static_assert(
        std::is_same_v<decltype(view2), std::ranges::owning_view<std::vector<int>>>,
        "Expected owning_view for rvalue"
    );
    return 0;
}
```

rewrite ex1 using ranges.
ex4.cpp

An example: Custom view and adaptor

Create a moving average algorithm. Given a range of numbers A. Produce a new range B with the moving average of the numbers in A.

need some kind of stateful iterator, not provided by std

```
int main() {  
    std::vector<double> data = {1, 2, 3, 4, 5, 6};  
    auto avg3 = data | moving_average(3);  
    return 0;  
}
```

Define a struct 'moving_average'

Concept used to constrain templated parameter r

This is the adaptor, the view factory.

```
struct moving_average {  
    std::size_t window_;  
    friend auto operator|(std::ranges::viewable_range auto&& r, moving_average const& self) {  
        return moving_average_view(std::forward<decltype(r)>(r), self.window_);  
    }  
};
```

```

template <std::ranges::input_range V>
requires std::ranges::view<V> && std::is_arithmetic_v<std::ranges::range_value_t<V>>
class moving_average_view
    : public std::ranges::view_interface<moving_average_view<V>>
{
    V base_;
    std::size_t window_;
    using T = std::ranges::range_value_t<V>;

public:
    template <std::ranges::viewable_range R>
    moving_average_view(R&& r, std::size_t w)
        : base_(std::forward<R>(r)), window_(w) {}
    struct iterator {
        ...
    };
    iterator begin() {
        return iterator(std::ranges::begin(base_),
                        std::ranges::end(base_),
                        window_);
    }
    std::default_sentinel_t end() const noexcept {
        return {};
    }
};

```

```

struct iterator {
    using value_type = T;
    using difference_type = std::ptrdiff_t;
    using iterator_category = std::input_iterator_tag;
    std::ranges::iterator_t<V> it_;      // current source iterator
    std::ranges::sentinel_t<V> end_;    // source range end
    std::deque<T> buf_;
    std::size_t window_;
    iterator() = default;
    iterator(std::ranges::iterator_t<V> it,
            std::ranges::sentinel_t<V> end,
            std::size_t window)
        : it_(it), end_(end), window_(window) {}
    value_type operator*() const {
        // produce a moving average from buf_
    }
    iterator& operator++() {
        ++it_;          // advance the source iterator
        // update buf_ with the latest value
        return *this;
    }
    friend bool operator==(const iterator& i, std::default_sentinel_t) {
        return i.it_ == i.end_; // the iterator ends when it reaches the source sentinel
    }
    friend bool operator==(std::default_sentinel_t, const iterator& i) {
        return i == std::default_sentinel_t{};
    }
}

```

The key idea is really simple:

- `it_` is the source iterator
- `end_` is the source sentinel
- the view's `end()` is the default sentinel of the view
- the relation `it_ == end_` decides when the custom iterator has exhausted the source range

The power of `std::views::all_t`

CTAD rule, deduce parameter V from parameter R.

```
template <class R>
moving_average_view(R&&, std::size_t)
-> moving_average_view<std::views::all_t<R>>>;
```

`std::views::all_t` deduces to `ref_view` or `owning_view` depending on R.

```
int main() {
    std::vector<double> data = {1, 2, 3, 4, 5, 6};
    auto avg3 = data | moving_average(3);
    static_assert(std::same_as<decltype(avg3), moving_average_view<std::ranges::ref_view<std::vector<double>>>>);
    auto avg4 = std::move(data) | moving_average(3);
    static_assert(std::same_as<decltype(avg4), moving_average_view<std::ranges::owning_view<std::vector<double>>>>);
    return 0;
}
```

Adapter closure now only works in pipeline

Free pipeline operator --> inherit from `std::ranges::range_adaptor_closure`

```
struct moving_average : std::ranges::range_adaptor_closure<moving_average>
{
    std::size_t window_;
    auto operator()(std::ranges::viewable_range auto&& r) const {
        return moving_average_view(std::forward<decltype(r)>(r), window_);
    }
};

int main() {
    std::vector<double> data = {1,2,3,4,5,6};
    auto avg_pipe = data | moving_average{.window_ = 3};
    static_assert(std::same_as<decltype(avg_pipe), moving_average_view<std::ranges::ref_view<std::vector<double>>>>);
    moving_average ma{.window_ = 3};           // closure object
    auto avg3 = ma(data);                       // call operator() directly
    static_assert(std::same_as<decltype(avg3), moving_average_view<std::ranges::ref_view<std::vector<double>>>>);
}
```

ex5.cpp

C++ Ranges have a "Type-Safe Tombstone"

```
auto get_data() { return std::vector{1, 2, 3}; }
int main()
{
    // std::ranges::find returns std::ranges::dangling.
    auto it = std::ranges::find(get_data(), 2);
    return *it;
}
```

Runtime UB error intercepted by the compiler:

```
<source>: In function 'int main()':  
<source>:9:12: error: no match for 'operator*' (operand type is 'std::ranges::dangling')  
    9 |         return *it;  
      |                ^~~  
Compiler returned: 1
```

Why does `std::find` return dangling while `views::all` returns an `owning_view`?

- Algorithms return iterators. They cannot own the container.
- Views are objects. When given an Rvalue, they move the data into an `owning_view`.

```
//moving_average_view<std::ranges::owning_view<std::vector<int>>>  
auto v = get_data() | std::views::filter(is_even);  
// 'it' cannot hold the vector. To prevent a crash, it becomes 'dangling'.  
auto it = std::ranges::find(get_data(), 2);
```

Borrowed ranges

Observation: Not all r-value ranges are unsafe

If the range dies, the iterators remain valid

```
auto it = std::ranges::find(std::string_view{"Hello"}, 'e');  
static_assert(!std::same_as<decltype(it), std::ranges::dangling>);
```

```
template<typename T>
struct MyCursor {
    T* ptr;
    T* begin() const { return ptr; }
    T* end() const { return ptr + 10; }
};

template<typename T>
inline constexpr bool std::ranges::enable_borrowed_range<MyCursor<T>> = true;

int main() {
    auto it = std::ranges::find(MyCursor<int>{nullptr}, 42);
    return *it;
}
```

```
template<std::ranges::input_range R>
decltype(auto) get_first(R&& r) {
    if constexpr (std::ranges::borrowed_range<R>) {
        return *std::ranges::begin(r); // Returns Reference (Fast)
    } else {
        using V = std::ranges::range_value_t<R>;
        return static_cast<V>(*std::ranges::begin(r)); // Returns Copy (Safe)
    }
}

int main() {
    decltype(auto) val = get_first(std::vector{1, 2});
}
```

Custom sentinel: keeps iterating until std::optional is empty

```
struct element_iterator {
    explicit element_iterator(const char* p)
        : cur_(p) {
        current_ = read_next();
    }
    const Element& operator*() const { return *current_; }
    element_iterator& operator++() {
        current_ = read_next();
        return *this;
    }
    friend bool operator==(const element_iterator& it, std::default_sentinel_t) {
        return !it.current_.has_value();
    }
private:
    const char* cur_ = nullptr;
    std::optional<Element> current_;
    std::optional<Element> read_next() { }
};
```

std::ranges::view_interface

Avoiding boilerplate

.empty(), .size(), .front(), .back

```
class element_range : public std::ranges::view_interface<element_range> {
public:
    explicit element_range(const char* text) : text_(text) {}
    element_iterator begin() const { return element_iterator{text_}; }
    std::default_sentinel_t end() const { return {}; }
private:
    const char* text_;
};
```

Compile-time

- Heavy lifting is done by the type system.
- building the pipeline is done at compile time
- runtime; iterating over the items.
- You could get many specializations for a view (moving_average_view)
- code size can explode
- compile times can explode, any_view: P3411R0

William TODO: zip issue bij tomra

ex6.cpp

ex7.cpp

ex8.cpp

We bring **high-tech** to life